School of Electrical and Information Engineering
University of the Witwatersrand, Johannesburg
ELEN4020 — Data Intensive Computing in Data Science

Laboratory Exercise 2

1 Outcome

This is a group laboratory exercise. The main purpose of this laboratory exercise is to learn to use shared memory programming libraries, PThread and OpenMP, for simple scientific applications of manipulating very large matrices that are maintained in memory. The main outcome will be:

1. learning to use the PThread and OpenMP parallel programming libraries.

2. learning to compile and run parallel shared memory programs in a Linux environment.

3. learning to evaluate the performance of your parallel programs by timing.

2 Problem Description

You are expected to write a program that computes the transpose of a square matrix $A[N_0][N_1]$ in-place for $N_0 = N_1$ and for $N_0 = 128$, 1024, 2048 and 4096. Computing the transpose in place requires that you do not create a second $N_0 \times N_1$ matrix, say $A^T[N_0][N_1]$, and copy elements of the original matrix A, into A^T to contain the transposed elements of the original matrix. The elements of the original matrix should simply be reshuffled, using a small temporal storage space much much less than $N_0 \times N_1$, to generate the transposed elements in the original matrix.

For example if the original matrix is a 4×4 matrix $A[4][4]$ as shown below,

$$A[4][4] = \begin{bmatrix} a & b & c & d \\ e & f & g & h \\ i & j & k & l \\ m & n & o & p \end{bmatrix}$$

the transposed matrix should be

$$A^T[4][4] = \begin{bmatrix} a & e & i & m \\ b & f & j & n \\ c & g & k & o \\ d & h & l & p \end{bmatrix}$$

The number of threads in both the PThread program and the OpenMP program should be chosen appropriately according to the suggested algorithm explained below. For each of the values of N_0 and number of threads used, evaluate the time it takes to perform the transposition of the matrix, not including the time to populate the original matrix. Elements of the matrix $A[N_0][N_1]$ should be a 4-byte integer, generated as random values between 0 and N_0.

For the use of PThreads two algorithms will be considered:

1. A Diagonal-Threading algorithm.

2. A Block-Oriented-Threading algorithm.

For OpenMP:

1. A Naive-Threading algorithm.

2. A Diagonal-Threading algorithm.

3. A Block-Oriented-Threading algorithm.

2.1 OpenMP Naive-Threaded Algorithm

A simple non-threaded approach, or serial algorithm, is to transpose the given matrix with a simple nested loop by swapping the elements $A\langle i, j\rangle$ with $A\langle j, i\rangle$. Call this method the Naive-Algorithm.

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24
25	26	27	28	29	30	31	32
33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48
49	50	51	52	53	54	55	56
57	58	59	60	61	62	63	64

1	9	17	25	33	41	49	57
2	10	18	26	34	42	50	58
3	11	19	27	35	43	51	59
4	12	20	28	36	44	52	60
5	13	21	29	37	45	53	61
6	14	22	30	38	46	54	62
7	15	23	31	39	47	55	63
8	16	24	32	40	48	56	64

Figure 1: Example matrix (left) and transposed matrix (right)

A naive OpenMP threading will instrument the serial algorithm by inserting `#PRAGMA` before for loops, etc., to automatically parallelize the code if compiled as an OpenMP application.

2.2 Diagonal-Threading Algorithm

In this algorithm, threads are generated for each diagonal elements. For each diagonal position, the corresponding row elements to the right and the column elements below are interchanged by the threads assigned to the position (See Figure 2a).

2.3 Block-Oriented-Threading Algorithm

The Block-Oriented-Threading algorithm can be perceived as 2-levels of transpositions. The matrix is chunked into blocks. The elements within each block are first transposed and then the blocks are then transposed to generate the overall transposed matrix. See the illustration of the algorithm below in Figure 2b.

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24
25	26	27	28	29	30	31	32
33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48
49	50	51	52	53	54	55	56
57	58	59	60	61	62	63	64

(a) Diagonal-Threading

Block-Oriented-Threading (columns 0, 1, 2, 3; rows 0, 1, 2, 3):

	0	1	2	3
0	1 9 / 2 10	A_{10}	33 41 / 34 42	A_{30}
1	A_{01}	19 27 / 20 28	A_{21}	51 59 / 52 60
2	5 13 / 6 14	A_{12}	37 45 / 38 46	A_{32}
3	A_{03}	23 31 / 24 32	A_{23}	55 63 / 56 64

(b) Block-Oriented-Threading

Figure 2: Two parallel algorithms to transpose a matrix

3 Work Schedule

The work will involve:

- developing and implementing a multi-threaded algorithm for computing the transpose of a matrix in place in PThread and OpenMP. Your algorithm should be as efficient and also scalable as much as possible.

- generating a comparative table of the performance of your algorithm as in the sample table below.

$N_0 = N_1$	Naive	PThreads		OpenMP		
		Diagonal	Blocked	Naive	Diagonal	Blocked
128						
1024						
2048						
4096						

4 Deliverables

- A lab-report analysing the performance of the algorithms as N_0 changes. Max 2 pages and one report per group.
  - Provide a short description of your algorithm to the problem of 2-D matrix transposition and also the pseudo-code of your routines.

- Your code for marking in your initialised group `git` repository in the Github Organisation.
  - Ensure your repository has an appropriate `README` describing the build process.